


ML Presentation

Kenneth Hansen

Mads Helt

Mike Chuang

Ilyas Shomayev




Summary:

- EDA
- Pre-processing
 - Imputing Missing Values
 - Encoding Categorical Features
 - Skewness
 - Feature Engineering
- Modeling
 - Model Performance
 - Model Tuning
 - Model Ensembling
- Results
- RMSLE score using CV on training set
- Kaggle score



Pre-Processing: Missing Values



There is a substantial amount of missing data. We differentiate columns on a categorical/numeric basis.

Pre-Processing: Missing Values

- Categorical Features:
 - Impute Missing with “None”
 - Imputed with mode of feature
 - MSZoning, Functional, Exterior1st

Categorical features

```
list_none = ['PoolQC', 'GarageType', 'GarageFinish', 'GarageCond', 'GarageQual', 'BsmtQual', 'BsmtCond', 'KitchenQual',
            'BsmtExposure', 'BsmtFinType2', 'BsmtFinType1', 'MasVnrType', 'Alley', 'Fence', 'FireplaceQu', 'MiscFeature']

# Replace with "None"
df_all[list_none] = df_all[list_none].fillna('None')

# replace with the mode
df_all['MSZoning'] = df_all['MSZoning'].fillna(df_all['MSZoning'].mode()[0])
df_all['Functional'] = df_all['Functional'].fillna(df_all['Functional'].mode()[0])
df_all['Exterior1st'] = df_all['Exterior1st'].fillna(df_all['Exterior1st'].mode()[0])
```

- Numeric Features:
 - Impute Missing with “0”
 - Impute with Median of feature
 - LotFrontage

Numeric features

```
list_null = ['GarageYrBlt', 'GarageArea', 'GarageCars', 'BsmtFinF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
            'BsmtFullBath', 'BsmtHalfBath', 'MasVnrArea', 'Exterior2nd', 'SaleType', 'Electrical']

# Replace with 0
df_all[list_null] = df_all[list_null].fillna(0)

# Replace with the median of the neighborhood
df_all['LotFrontage'] = df_all.groupby('Neighborhood')['LotFrontage'].transform(lambda x: x.fillna(x.median()))
df_all.loc[df_all.LotFrontage.isnull(), 'LotFrontage'] = np.sqrt(df_all.loc[df_all.LotFrontage.isnull()].LotArea)
```


Pre-Processing: Encoding

Ordinal Categorical Features: Converted to numeric based on scale:

```
qual_dict = {"None": 0, "Po": 1, "Fa": 2, "TA": 3, "Cd": 4, "Ex": 5}
df_all["ExterCond"] = df_all["ExterCond"].map(qual_dict).astype(int)
df_all["ExterQual"] = df_all["ExterQual"].map(qual_dict).astype(int)
df_all["BmntQual"] = df_all["BmntQual"].map(qual_dict).astype(int)
df_all["BmntCond"] = df_all["BmntCond"].map(qual_dict).astype(int)
df_all["HeatingQC"] = df_all["HeatingQC"].map(qual_dict).astype(int)
df_all["KitchenQual"] = df_all["KitchenQual"].map(qual_dict).astype(int)
df_all["GarageQual"] = df_all["GarageQual"].map(qual_dict).astype(int)
df_all["GarageCond"] = df_all["GarageCond"].map(qual_dict).astype(int)
df_all["FireplaceQu"] = df_all["FireplaceQu"].map(qual_dict).astype(int)
```

9 other variables received this treatment but had different scales.

Inspecting the boxplot distribution for 'FireplaceQu' shows that retaining ordinality is important.



Pre-Processing: Encoding

Non-ordinal Categorical Features: :

```
df_all = df_all.replace({"MSSubClass": {20: "SC10", 30: "SC30", 40: "SC40", 50: "SC50", 60: "SC60", 70: "SC70", 80: "SC80", 90: "SC90", 100: "SC100", 120: "SC120", 150: "SC150", 160: "SC160"}, "MoSold": {1: "Jan", 2: "Feb", 3: "Mar", 4: "Apr", 5: "May", 6: "Jun", 7: "Jul", 8: "Aug", 9: "Sep", 10: "Oct", 11: "Nov", 12: "Dec"}  
df_all["YrSold"] = df_all["YrSold"].astype(str)
```

MSSubClass, MoSold, and YrSold were initially numeric features but actually hold categorical data

```
# Label encoding non-order categorical features  
from sklearn.preprocessing import LabelEncoder  
columns = ('PavedDrive', 'Alley', 'Street', 'CentralAir', 'MSSubClass')  
for column in columns:  
    lbl = LabelEncoder()  
    lbl.fit(list(df_all[column].values))  
    df_all[column] = lbl.transform(list(df_all[column].values))  
print(df_all.shape)
```

(2917, 80)


```
df_all = pd.get_dummies(df_all, drop_first = True)  
print(df_all.shape)
```

(2917, 237)

The LabelEncoder function performs one-hot encoding within the listed columns

Performing one-hot encoding on the remaining dataset then yields additional columns, resulting in 237 total columns.

Pre-Processing: Skewness



We discovered the Target Variable (SalePrice) was skewed, so we perform a log transformation to normalize its distribution. We then investigated the skewness of the predictor variables to understand if further transformation was necessary.

Pre-Processing: Skewness

| Variable Name | Positive Skewness | Variable Name | Negative Skewness |
|---------------|-------------------|---------------|-------------------|
| Fireplaces | 0.725277917 | Street | -15.49475602 |
| TotRmsAbvGrd | 0.74923209 | LandSlope | -4.973253614 |
| ExterGrade | 0.782428205 | Functional | -4.961674871 |
| ExterQual | 0.783456143 | GarageYrBlt | -3.904632328 |
| 2ndFlrSF | 0.861555523 | BsmtCond | -3.602661074 |
| BsmtUnfSF | 0.919688213 | CentralAir | -3.45755483 |
| AvgRoomSize | 0.931703531 | GarageCond | -3.381673401 |
| BsmtFinSF1 | 0.980644589 | GarageQual | -3.262259968 |
| TotLivArea | 1.009156621 | PavedDrive | -2.977741053 |
| GrLivArea | 1.06875039 | BsmtQual | -1.271610943 |
| LotFrontage | 1.103038596 | IndStope | -1.247972934 |
| BsmtExposure | 1.119066336 | | |
| 1stFlrSF | 1.257285977 | | |
| ExterCond | 1.315069293 | | |
| Fence | 1.753731433 | | |
| WoodDeckSF | 1.844791628 | | |
| OverallScore | 1.907677233 | | |
| AllPorchSF | 2.244499743 | | |
| OpenPorchSF | 2.529358203 | | |
| MasVnrArea | 2.621719301 | | |
| BsmtFinType2 | 3.150951371 | | |
| BsmtHalfBath | 3.929995969 | | |
| ScreenPorch | 3.945101226 | | |
| EnclosedPorch | 4.002344092 | | |
| BsmtFinSF2 | 4.14450336 | | |
| KitchenAbvGr | 4.300550114 | | |
| BsmtFin1 | 5.27207993 | | |
| LowQualFinSF | 5.29045605 | | |
| UnfFlrSF | 13.4491969 | | |
| PoolArea | 17.08880449 | | |
| PoolQC | 20.84421041 | | |
| PoolFence | 21.93901217 | | |

We also found 6 features with skewness <0.3. However, upon further inspection we found that only dropping the “Utilities” variable would result in a minimal loss of predictive power.

We found 43 features to exhibit skewness > 0.7. Therefore, we performed a boxcox transformation to normalize each of them.

| Variables with <0.3 skewness |
|------------------------------|
| Condition2 |
| Heating |
| PoolQC |
| RoofMatl |
| Street |
| Utilities |

Pre-Processing: Feature Engineering

Feature Engineering: Created 11 new variables based off of feature interactions

```
# 1: Total square feet
df_all['TotLivArea'] = df_all.TotalBsmSF + df_all['1stFlrSF'] + df_all['2ndFlrSF']

# 2: # Total number of bathrooms
df_all['totalBath'] = df_all['BsmFullBath'] + (0.5 * df_all['BsmHalfBath']) + \
df_all['FullBath'] + (0.5 * df_all['HalfBath'])

# 3: BsmUnFinRatio
df_all['BsmUnFinRatio'] = df_all.BsmUnFinSF / df_all.TotalBsmSF
df_all['BsmUnFinRatio'] = df_all['BsmUnFinRatio'].fillna(0)

# 4: AreaPerCar
df_all['AreaPerCar'] = df_all.GarageArea / df_all.GarageCars
df_all['AreaPerCar'] = df_all['AreaPerCar'].fillna(0)

# 5: AvgRoomSize
df_all['AvgRoomSize'] = df_all.GCLivArea / df_all.TotalRoomsABVGRD

# 6: GarageScore
df_all['GarageScore'] = df_all.GarageCond * df_all.GarageType

# 7: OverallGrade
df_all['OverallGrade'] = df_all["OverallQual"] * df_all["OverallCond"]

# 8: OverallScore
df_all['OverallScore'] = df_all["OverallGrade"] * df_all['TotLivArea']


#9: Exterior Grae
df_all['ExterGrade'] = df_all["ExterQual"] * df_all["ExterCond"]

#10: All Porch SF
df_all['AllPorchSF'] = df_all["OpenPorchSF"] + df_all["EnclosedPorch"] + \
df_all["3SeasonPorch"] + df_all["ScreenPorch"]

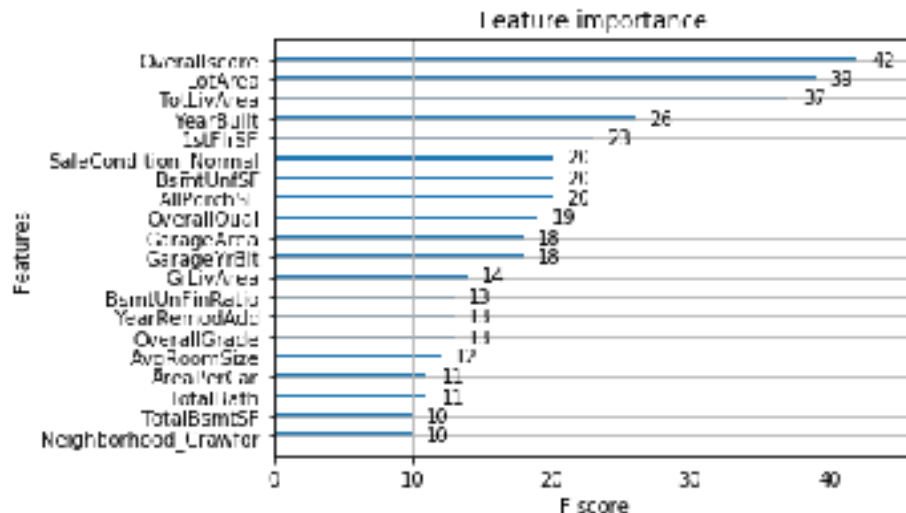
#11: Age
df_all['Age'] = 2010 - df_all['YearBuilt']
```

Age is the only variable that can be considered as a feature representation rather than an interaction. Room to explore creating further indicator/categorical features.

Model Feature Importance:



Lasso Coefficients (most important variables)



Xgboost – most important features

Parameter Tuning:

- Grid Search
- Bayesian Optimization




Modeling: Parameter Tuning

Linear Models: After running 5000 fits through GridSearchCV, optimal parameters are below:

```
Results for ridge:
Best parameters: {'alpha': 5.7746151523598144, 'random_state': 42}
Fitted MSE: 0.0126834105009
Fitted RMSE: 0.112620648644
-----
Results for lasso:
Best parameters: {'alpha': 0.00020923834803969769, 'random_state': 42}
Fitted MSE: 0.0123602338342
Fitted RMSE: 0.111176588516
-----
Results for elastic net:
Best parameters: {'alpha': 0.0043178797273320212, 'l1_ratio': 0, 'random_state': 42}
Fitted MSE: 0.0126836348432
Fitted RMSE: 0.112621644646
```


With an Elastic Net `l1_ratio` of 0, this means that that the ridge regression may yield the best results over a Lasso. The Lasso's near 0 alpha result also shows that a straightforward linear regression may also produce positive results.

Model Evaluation:



\hat{Y} VS. actual SalePrice (lasso model):

We observe a strong linear trend in the results



Residual Plot (lasso) \hat{Y} vs. Residuals

We do not observe any significant outliers.

Tuned Model Performances:

Ridge:

Mean RMSLE: 0.111579632589
Min RMSLE: 0.0956569591734
Max RMSLE: 0.137629970056
Std RMSLE: 0.0134479803353

Xgboost Alternative:

Mean RMSLE: 0.121502812757
Min RMSLE: 0.101831128608
Max RMSLE: 0.142983312984
Std RMSLE: 0.0155486014157

Lasso:

Mean RMSLE: 0.110223352332
Min RMSLE: 0.0960129858893
Max RMSLE: 0.13684855486
Std RMSLE: 0.0133668504869

Random Forest Aggressive:

Mean RMSLE: 0.113899715187
Min RMSLE: 0.096376429415
Max RMSLE: 0.133121484438
Std RMSLE: 0.0143629567948

Xgboost:

Mean RMSLE: 0.115406681482
Min RMSLE: 0.100213833837
Max RMSLE: 0.135809223636
Std RMSLE: 0.0137285687266

Random Forest Conservative:

Mean RMSLE: 0.132577372561
Min RMSLE: 0.111758666571
Max RMSLE: 0.158920830563
Std RMSLE: 0.014566172841

RESULTS

ridge, model_xgb, RForest: **0.1151**

ENet, lasso, ridge, model_xgb: **0.1086**

lasso, ridge, model_xgb: **0.1082**

lasso ridge, model_xgb, RForest: **0.1115**

ENet, lasso, ridge, model_xgb (all new features): **0.1087**

ENet, lasso, ridge, model_xgb (excluded last new variables): **0.1088**

ENet, lasso, ridge, model_xgb (excluded PoolQC and Street): **0.1089**

lasso, model_xgb (excluded PoolQC and Street): **0.1089**

lasso, model_xgb (all new features): **0.1085**

lasso, model_xgb (excluded low variance variables): **0.1087**

lasso, model_xgb (lower lambda for box-cox): **0.1086**

lasso, model_xgb (with two new derived features): **0.1084**

lasso, model_xgb (ex. levels not in test set): **0.1086** - best on Kaggle leaderbord (0.11763)

lasso, model_xgb (ex. levels not in test set and "Age"): **0.1088**

lasso, model_xgb, RForest_agg (ex. levels not in test set): **0.1117**


lasso, model_xgb, RForest_agg, RForest_con (ex. levels not in test set): **0.1155**

lasso, model_xgb, model_lgb: **0.1086**

lasso, model_lgb: **0.1076**

lasso, model_xgb_al (gridsearchCV parameters): **0.1108**

lasso, model_xgb (with new variable "SoldYr"): **0.1085**



Results:

0.11774 RMSLE ON PUBLIC TEST

Thank You!